

Toward a Sparse and Interpretable Audio Codec

John Vinyard

Austin TX, USA

john.vinyard@gmail.com

Abstract—Most widely-used modern audio codecs, such as Ogg Vorbis and MP3, as well as more recent “neural” codecs like Meta’s Encodec [1] or the Descript Audio Codec [2] are based on block-coding; audio is divided into overlapping, fixed-size “frames” which are then compressed. While they often yield excellent reproductions and can be used for downstream tasks such as text-to-audio, they do not produce an intuitive, directly-interpretable representation. In this work, we introduce a proof-of-concept audio encoder that represents audio as a sparse set of events and their times-of-occurrence. Rudimentary physics-based assumptions are used to model attack and the physical resonance of both the instrument being played and the room in which a performance occurs, hopefully encouraging a sparse, parsimonious, and easy-to-interpret representation.

1. INTRODUCTION

This work imagines a future audio codec where some types of musical composition could take place directly in the audio codec space. The text-to-audio paradigm works well for non-musicians creating background music for movies, advertisements or social media content, but it is our view that experienced musicians and composers work in a “space” not fully captured by language alone. We theorize that they will continue to prefer a finer-grained representation that provides interpretability and control at multiple scales.

This work does not seek to produce a generative model of musical audio, but could serve as the underlying encoding on which generative models are trained. It is the authors’ intuition that models trained on this quasi-symbolic representation might have a much more thorough “understanding” of the content they produce, given the point-cloud-like nature of the signal. Instead of predicting the next arbitrarily-sized frame, the generative model would be predicting the relationships between “events”. Long-term coherence in musical generation continues to be a challenge, and we speculate that many models spend large shares of their capacity learning to reproduce physical resonance, reducing the share that may model the human forces that drive them in interesting, musical directions.

All model and experiment code is implemented using the PyTorch [3] Python library and can be found on GitHub¹. Audio examples can be heard at this [https URL](https://github.com/JohnVinyard/matching-pursuit)²

2. PRIOR WORK

We take inspiration from a wealth of previous works dealing with sparse and interpretable representations of audio signals. Matching pursuit [5] is an iterative algorithm that decomposes a signal into a sparse set of representative “atoms”. It produces a representation of audio signals similar to the one used in granular synthesis [6], which represents audio as point-cloud-like structure of audio “grains” arranged in time. While the notion of “events in time” feels intuitive to many of us, both approaches traditionally require hundreds or thousands of “atoms” or “grains” to accurately reconstruct a signal. These audio quanta are often orders of shorter than the time scale on which we typically conceive of audio events. Furthermore, matching




Fig. 1: In this visualization of the codec representation, we see that events can overlap and vary in length. Time is along the x-axis, event positions are along the y-axis and event colors are chosen by applying t-SNE [4] to the set of 32 event vectors, targeting a single dimension for the y-axis and three dimensions to represent an RGB color value.

pursuit, at least when operating naively in the time domain can spend capacity on perceptually irrelevant details and struggle to reproduce noisy signals in a convincing way. In recent years, attempts have been made to use longer atoms and parameterize dictionaries of atoms in more meaningful ways [7]. In this work we hope to continue the trend toward fewer and more meaningful “atoms” or “events” that exist at scales native to a human musician.

Unsupervised audio source separation, such as in [8] seeks to separate mixtures of instruments or voices into distinct tracks, or “stems” as they are often referred to in music production, while unsupervised music transcription [9] seeks to infer a musical score, or set of instructions sufficient for a musician to reproduce the piece. Other recent works seek to produce meaningful latent representations by encouraging sparsity as part of the training objective [10]. In this work, we seek to infer a score of sorts, where each encoded note or event contains enough information to create a perceptually indistinguishable reproduction of the event.

Finally, source-excitation synthesis [11] models natural sounds as injections of energy into a system, often represented by a burst or pulse-train of white noise, and the resonance of that system in response to the event, often represented by one or more time-varying filters that emulate the resonant modes of the system along with any deformations of the system that occur as it resonates. In this work, we use this synthesis technique as the basis of the event decoder. Empirically, this strong inductive bias seems to aid in the promotion of sparsity (few events) and to provide us with intermediate states that are open to inspection and interpretation.

3. AUDIO CODEC DETAILS

Our proposed audio encoding consists of a sequence of two-tuples, each consisting of a scalar event time and the parameters required for the decoder to render the event, transforming it into “raw” PCM audio samples. In addition to straightforward seeking and slicing,

¹<https://github.com/JohnVinyard/matching-pursuit>

²<https://blog.cochlea.xyz/sparse-interpretable-audio-codec-paper.html>

this encoding also makes it possible to filter a subset of events in a particular audio slice according to some criteria, an operation that is less straightforward in a block-coding scheme.

3.1. Compression Rate

While our main aim in this work is a representation that is sparse, interpretable and easy-to-manipulate, a compressive representation is one important aspect of a successful codec and should be discussed. For our experiment, the encoder is run for a fixed number of steps, 32 in our case, and produces a two-tuple of time-of-occurrence and 32-dimensional event vector. Each event time can be stored as a single scalar value. Because we encode 2^{16} samples, or around 2.98 seconds of audio at a time, at a 22050hz sampling rate, we arrive at:

$$\frac{2^{16} \text{ samples}}{(32 * 32) + 32} \approx 62x \quad (1)$$

4. MODEL

4.1. Audio Representation

In this work, we seek an audio representation that can be iteratively decomposed, without wasting capacity on perceptually-irrelevant details. As a concrete example, removing energy from a time-domain representation of noise can be very difficult, as the only way to do so is to match the signal *exactly* even though small time and/or frequency shifts would be imperceptible to the common listener.

We choose the widely-used STFT magnitude spectrogram representation of audio, with a window size of 2048 and a hop size of 256. The 75% overlap means that we can discard phase but still recover *most* perceptually relevant details, including phase relationships.

While the model only encodes segments of 2^{16} samples at a time, it *analyzes* segments of 2^{17} , encoding and removing only events that begin in the first half of the audio segment. This enables us to implement a streaming encoder. The model's input at each step is an STFT magnitude spectrogram with 1025 real-valued coefficients and 512 "frames" in the time dimension.

4.2. Encoder

In this work, the encoder is parameterized as an anti-causal, dilated convolutional network with a kernel size of 2 at each layer and dilation sizes of [1, 2, 4, 8, 16, 32, 64].

In our experiments, the encoder runs for a fixed number of steps (32). At each step, it transforms the residual spectrogram into a tensor with 32 channels and 512 frames, the same time dimension as the input spectrogram. The "residual" spectrogram at the first step is simply the STFT magnitude spectrogram computed from the input audio.

The encoder then selects a *single* event vector along the time dimension, setting all other 32-dimensional vectors to zero. When we take the norm of each position along the time dimension, we derive a one-hot vector that will be used for the coarse-grained timing of the event when it is "scheduled" after the event decoder has generated the event.

Once generated and scheduled, the STFT magnitude spectrogram of the newly-generated event is computed and subtracted from the residual spectrogram, preparing the model to begin the next iteration.

While our initial experiments use a fixed number of encoding steps, we believe that we will be able to devise more intelligent stopping conditions in future versions of this work that take advantage of the variable information density across different audio segments. One simple and obvious possibility is to choose a small norm as the threshold at which the sound would no longer be audible to listeners.




Fig. 2: High-level model architecture. The encoder and decoder work together to incrementally remove energy from the input representation, an STFT-based magnitude spectrogram. At each step, the encoder produces a single event vector and time-of-occurrence. The decoder generates an audio event, positions it in time, and subtracts it from the input representation.




Fig. 3: Decoder Block

4.3. Streaming Algorithm

To enable a streaming encoding algorithm, the encoder masks the second-half of its output just before choosing the next event vector to be decoded and removed. This means that the encoder is always choosing events that begin in the first-half of its input, but the events it chooses may into the second half.

To enforce this focus on the first half of the analyzed signal, we mask encoder output before choosing an event, *and* multiply the second half of the signal by a linear gradient beginning at 1 and ending at 0 which extends from sample 2^{16} to sample 2^{17} . This means that model is penalized less for failures to remove energy from the second half of the signal and that it does not attempt to produce overly-long events that stretch far into the second-half.

4.4. Decoder

Instead of a more traditional convolutional upsampling network, we choose a source-excitation model for the event decoder in the hope that it encourages both sparser and easier-to-interpret events.

In Figure 3, we can see a diagram of a single decoder block. Each block has a few distinct parts with clear interpretations. Inputs to each block are generated by a small, “multi-head” MLP whose input is the original 32-dimensional event vector.

The input to each block is a source signal, which takes the form of one or more bursts of noise that represent the injection of energy into the system. The noise burst is convolved with a number of decaying resonances, represented in our work as relatively long FIR filters composed of exponentially decaying sinusoids. Each convolved resonance is then multiplied by a time-varying mask and summed, emulating a filter with time-varying parameters, or in a more physical interpretation, the deformations applied over time to a resonating object.

We refer to the number of resonances chosen for a particular block as its expressivity. Some simple physical resonances, such as a tuning fork, might only require that expressivity=1, while others, such as a vibrato violin note, might require a larger expressivity number.

Finally a two-channel gain is applied to the original source and the time-varying, filtered-signal, akin to a weighted skip connection.

These decoder blocks can be stacked, with the output of one block serving as the input or “impulse” fed to the next.

For this set of experiments, we choose three decoder blocks, with the last layer being fixed and given an expressivity value of 1. This block is initialized with a set of freely-available room impulse responses [12] that are commonly used to produce reverb effects. This makes it *possible* to disentangle instrument and room resonances.

For coarse-grained scheduling, the audio generated by the decoder is first convolved with the one-hot vector produced by the encoder. Fine-grained, sample-level scheduling is achieved using a frequency-domain time-shift, represented by a single scalar value, that is generated from the event vector by one of the decoder’s MLP “heads”.

4.5. Model Size

The model used in our experiments has a total of 45.1M parameters and occupies 200MB when serialized to disk. We feel that even this small model shows promise and that further experiments with larger models are warranted.

5. EXPERIMENT

Our primary experimental goal is to show that we can encode a large set of “natural” (non-synthesized) musical audio with diverse instruments and recording conditions into a representation that enjoys sparsity, interpretability and good reproduction quality.

5.1. Data Set

We train our model on the MusicNet dataset [13], an open dataset containing 33 hours of classical music, recorded in diverse spaces using a range of recording equipment of varying quality. The recordings often include leading silence, trailing applause, and incidental human sounds throughout (coughs, movement, etc).

While the dataset as a whole includes fine-grained score information about each musical piece, we do not use this component of the dataset, using *only* the audio signals to learn our unsupervised model.

5.2. Training Algorithm

During training, we repeatedly draw segments of audio that are approximately 5.94 seconds in length from the Music Net dataset at random, with length 2^{17} samples at 22050hz sampling rate and a batch size of two.

We train using a single NVIDIA GeForce RTX 3060 for approximately 76 hours.

Each batch is passed through the encoding and decoding process for a fixed number of steps, 32, in our case, and then an iterative loss is computed, encouraging each step to have removed as much energy (expressed as the L1 norm) from the signal as possible. This loss is minimized using the Adam optimizer with a learning rate of $1e^{-4}$. Empirically, we found that in early stages of training, the model tended to simply output silence as a pathological local minima. Empirically, the relatively small learning rate seemed to help move beyond these early-stage pathologies, but a learning-rate schedule that increases once the model has stabilized might yield better results overall.

5.3. Loss Functions

The loss function is also iterative and greedy. Using the same magnitude spectrogram representation analyzed by the encoder, the loss function attempts to *maximize* the reduction in the L1 norm of the spectrogram at each step. Prior to computing the loss, the input events are sorted in order of descending L1 norm, such that the event with the most energy is subtracted first and the event with the least energy is subtracted last.

6. INTERPRETABLE AND MANIPULATABLE REPRESENTATION

In this section, we discuss the multiple scales of interpretability and manipulability provided by the proposed codec. We intuit that these properties will be appealing to musicians and sound designers, and speculate that these features may prove just as important as compression rates under some circumstances.

6.1. Events and Times-of-Occurrence

As shown in Figure 1, relatively few events are required to reconstruct musical audio segments. Information density and temporal structure is apparent at-a-glance. At this scale, it is possible to preview individual events, translate them in time, and add or remove events based on some criteria.

6.2. Events Vectors

Aside from a coarse-grained event time, the low-dimensional event vectors encode all information required to produce an audio event. This work does not explore the latent space learned by the encoder in-depth, but an interactive two-dimensional map of event vectors (see Figure 4) demonstrates that a nearest-neighbors exploration can locate interesting variations on a particular event.

6.3. Decoder Interpretation

While the high-level model architecture is decoder-agnostic, our choice of a source-excitation-inspired decoder in this first experiment means that further interpretation and manipulation is possible by inspecting intermediate decoder states. In Figure 5 we see spectrograms of the initial impulse, the resonance, and finally the selected room impulse-response for a single event. While this work does not include an in-depth exploration of decoder interpretation and manipulation, it’s possible that different impulses, resonances, or room impulse responses could be chosen, while other aspects are held constant in order to subtly or profoundly influence the rendering of a particular event.




Fig. 4: This scatterplot shows event vectors from a large number of audio segments mapped onto a 2D-plane using t-SNE [4]. Exploring nearby neighbors can locate variations of a query event. An interactive version of this scatterplot can be explored [here](#).




Fig. 5: Here we can see intermediate stages of the source-excitation-based decoder at work. From left to right, we see A. a spectrogram of the initial noisy impulse B. a spectrogram of the noisy impulse convolved with the chosen decaying resonance and C. the result from step B. convolved with a room impulse response.

7. CONCLUSION

In this work, we propose an audio codec optimized for interpretability and ease-of-manipulation and then discuss a simple, small-scale reference implementation of an encoder and decoder network. We find that while subjective reproduction quality falls short in this iteration, the results achieved with a relatively small network are encouraging, and the properties of the codec that promote intuitive understanding of the audio content are worth further study.

8. FUTURE WORK

Given the encouraging small-scale results, we feel that there are several future experiments ripe for exploration.

8.1. Perceptually-Inspired Losses

Our sense is that perceptual audio losses are an under-explored area, and that adversarial losses are often introduced to account for a fundamental mismatch between what commonly-used audio representations measure and what is perceptually meaningful to humans. While the magnitude spectrogram representation used in this work may largely solve the problem of dealing with band-limited noise, it does not address higher-level and more complex invariances in human auditory perception, as discussed in [14]. Leveraging perceptual invariances in “textures”, such as background hiss and noise, will likely result in simpler encodings. Audio losses that are more perceptually-informed will undoubtedly yield more parsimonious encodings, as more model capacity will be spent on perceptually-relevant details.

8.2. Encoder Variants

The anti-causal dilated convolutional network used in this iteration, while simple and parameter-efficient, may be suboptimal as an encoder. UNet and transformer encoder architectures are both worth exploring in future versions of this research.

8.3. Decoder Variants

Anecdotally, we have observed that the roughly-physics-based inductive biases of the decoder encourage a sparser representation and better reconstruction quality, but many other decoder architectures are possible, with much to be learned from the larger field of differentiable digital signal processing [15].

Another possibility is that events could be routed to specialized decoders, akin to the switch transformer architecture [16].

8.4. Redundant Events and Sparsity

Finally, aside from the physics-based inductive bias in the decoder, this work makes no effort to further impose sparsity or penalize overly “verbose” representations. It is our observation that this iteration of the model frequently produces redundant, duplicative events that could be collapsed further.

Future work should explore sparsity and/or energy penalties, seeking to maintain high reproduction quality while producing the sparsest, or lowest-energy representation that can adequately explain the input signal.

REFERENCES

- [1] A. Défossez, J. Copet, G. Synnaeve, and Y. Adi, “High fidelity neural audio compression,” 2022. [Online]. Available: <https://arxiv.org/abs/2210.13438>
- [2] R. Kumar, P. Seetharaman, A. Luebs, I. Kumar, and K. Kumar, “High-fidelity audio compression with improved rvqgan,” 2023. [Online]. Available: <https://arxiv.org/abs/2306.06546>
- [3] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein, L. Antiga, A. Desmaison, A. Kopf, E. Yang, Z. DeVito, M. Raison, A. Tejani, S. Chilamkurthy, B. Steiner, L. Fang, J. Bai, and S. Chintala, “Pytorch: An imperative style, high-performance deep learning library,” in *Advances in Neural Information Processing Systems* 32. Curran Associates, Inc., 2019, pp. 8024–8035. [Online]. Available: <http://papers.neurips.cc/paper/9015-pytorch-an-imperative-style-high-performance-deep-learning-library.pdf>
- [4] L. van der Maaten and G. Hinton, “Visualizing data using t-SNE,” *Journal of Machine Learning Research*, vol. 9, pp. 2579–2605, 2008. [Online]. Available: <http://www.jmlr.org/papers/v9/vandemaaten08a.html>
- [5] S. Mallat and Z. Zhang, “Matching pursuits with time-frequency dictionaries,” USA, Tech. Rep., 1993.
- [6] I. Xenakis, *Formalized Music: Thought and Mathematics in Composition*. Bloomington: Indiana University Press, 1971.
- [7] J. Alfonso, “Reds: A new asymmetric atom for sparse audio decomposition,” *DAFx17*, 2017. [Online]. Available: https://www.dafx17.eca.ed.ac.uk/papers/DAFx17_paper_66.pdf
- [8] K. Schulze-Forster, G. Richard, L. Kelley, C. S. J. Doire, and R. Badeau, “Unsupervised music source separation using differentiable parametric source models,” 2023. [Online]. Available: <https://arxiv.org/abs/2201.09592>
- [9] T. Berg-Kirkpatrick, J. Andreas, and D. Klein, “Unsupervised transcription of piano music,” in *Advances in Neural Information Processing Systems*, Z. Ghahramani, M. Welling, C. Cortes, N. Lawrence, and K. Weinberger, Eds., vol. 27. Curran Associates, Inc., 2014. [Online]. Available: https://proceedings.neurips.cc/paper_files/paper/2014/file/e51e118b28def0cabd5237b4e1fec499-Paper.pdf
- [10] M. Lisboa and G. Bellec, “Spiking music: Audio compression with event based auto-encoders,” 2024. [Online]. Available: <https://arxiv.org/abs/2402.01571>
- [11] G. Fant, *Acoustic Theory of Speech Production : With Calculations based on X-Ray Studies of Russian Articulations*. The Hague, Netherlands: Mouton, 1960.
- [12] “Voxengo-ir,” <https://oramics.github.io/sampled/IR/Voxengo/>, accessed: 2025-05-06. [Online]. Available: <https://oramics.github.io/sampled/IR/Voxengo/>
- [13] J. Thickstun, Z. Harchaoui, and S. M. Kakade, “Learning features of music from scratch,” in *International Conference on Learning Representations (ICLR)*, 2017.
- [14] J. McDermott, A. Oxenham, and E. Simoncelli, “Sound texture synthesis via filter statistics,” in *2009 IEEE Workshop on Applications of Signal Processing to Audio and Acoustics, WASPAA 2009*, ser. IEEE Workshop on Applications of Signal Processing to Audio and Acoustics, 2009, pp. 297–300, 2009 IEEE Workshop on Applications of Signal Processing to Audio and Acoustics, WASPAA 2009 ; Conference date: 18-10-2009 Through 21-10-2009.
- [15] J. Engel, L. Hantrakul, C. Gu, and A. Roberts, “Ddsp: Differentiable digital signal processing,” 2020. [Online]. Available: <https://arxiv.org/abs/2001.04643>
- [16] W. Fedus, B. Zoph, and N. Shazeer, “Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity,” 2022. [Online]. Available: <https://arxiv.org/abs/2101.03961>